

Homework 4

[Help Center](#)

The due date for this quiz is **Mon 11 Jan 2016 4:00 PM PST**.

☐ In accordance with the Coursera Honor Code, I (Juan Tzu Chuan) certify that the answers here are my own work.

Question 1

Neural Network and Deep Learning

A fully connected Neural Network has $L = 2$; $d^{(0)} = 5$, $d^{(1)} = 3$, $d^{(2)} = 1$. If only products of the form $w_{ij}^{(\ell)} x_i^{(\ell-1)}$, $w_{ij}^{(\ell+1)} \delta_j^{(\ell+1)}$, and $x_i^{(\ell-1)} \delta_j^{(\ell)}$ count as operations (even for $x_0^{(\ell-1)} = 1$), without counting anything else, which of the following is the total number of operations required in a single iteration of backpropagation (using SGD on one data point)?

- ☐ 43
- ☐ 53
- ☐ 47
- ☐ 59
- ☐ none of the other choices

Question 2

Consider a Neural Network without any bias terms $x_0^{(\ell)}$. Assume that the network contains $d^{(0)} = 10$ input units, 1 output unit, and 36 hidden units. The hidden units can be arranged in any number of layers $\ell = 1, \dots, L - 1$, and each layer is fully connected to the layer above it. What is the minimum possible number of weights that such a network can have?

- ☐ 44
- ☐ none of the other choices
- ☐ 46
- ☐ 43
- ☐ 45

Question 3

Following Question 2, what is the maximum possible number of weights that such a network can have?

- ☐ none of the other choices
- ☐ 500
- ☐ 520
- ☐ 490
- ☐ 510

Question 4

Autoencoder

Assume an autoencoder with $\tilde{d} = 1$. That is, the $d \times \tilde{d}$ weight matrix \mathbf{W} becomes a $d \times 1$ weight vector \mathbf{w} , and the linear autoencoder tries to minimize

$$E_{in}(\mathbf{w}) = \frac{1}{N} \sum_{n=1}^N \|\mathbf{x}_n - \mathbf{w}\mathbf{w}^T \mathbf{x}_n\|^2.$$

We can solve this problem with stochastic gradient descent by defining

$$\text{err}_n(\mathbf{w}) = \|\mathbf{x}_n - \mathbf{w}\mathbf{w}^T \mathbf{x}_n\|^2$$

and calculate $\nabla_{\mathbf{w}} \text{err}_n(\mathbf{w})$. What is $\nabla_{\mathbf{w}} \text{err}_n(\mathbf{w})$?

- ☐ $2(\mathbf{x}_n^T \mathbf{w})^2 \mathbf{w} + 2(\mathbf{x}_n^T \mathbf{w})(\mathbf{w}^T \mathbf{w}) \mathbf{x}_n - 4(\mathbf{x}_n^T \mathbf{w}) \mathbf{x}_n$
- ☐ none of the other choices
- ☐ $2(\mathbf{x}_n^T \mathbf{w})^2 \mathbf{x}_n + 2(\mathbf{x}_n^T \mathbf{w})(\mathbf{w}^T \mathbf{w}) \mathbf{w} - 4(\mathbf{x}_n^T \mathbf{w}) \mathbf{w}$
- ☐ $(4\mathbf{w} - 4)(\mathbf{x}_n^T \mathbf{x}_n)$
- ☐ $(4\mathbf{x}_n - 4)(\mathbf{w}^T \mathbf{w})$

Question 5

Following Question 4, assume that noise vectors ϵ_n are generated i.i.d. from a zero-mean, unit variance Gaussian distribution and added to \mathbf{x}_n to make $\tilde{\mathbf{x}}_n = \mathbf{x}_n + \epsilon_n$, a noisy version of \mathbf{x}_n . Then, the linear denoising autoencoder tries to minimize

$$E_{in}(\mathbf{w}) = \frac{1}{N} \sum_{n=1}^N \|\mathbf{x}_n - \mathbf{w}\mathbf{w}^T (\mathbf{x}_n + \epsilon_n)\|^2.$$

For any fixed \mathbf{w} , what is $\mathcal{E}(E_{in}(\mathbf{w}))$, where the expectation \mathcal{E} is taken over the noise generation process?

- ☐ $\frac{1}{N} \sum_{n=1}^N \|\mathbf{x}_n - \mathbf{w}\mathbf{w}^T \mathbf{x}_n\|^2 + \mathbf{w}^T \mathbf{w}$

- ☐ none of the other choices
- ☐ $\frac{1}{N} \sum_{n=1}^N \|\mathbf{x}_n - \mathbf{w}\mathbf{w}^T \mathbf{x}_n\|^2 + \frac{1}{d} \mathbf{w}^T \mathbf{w}$
- ☐ $\frac{1}{N} \sum_{n=1}^N \|\mathbf{x}_n - \mathbf{w}\mathbf{w}^T \mathbf{x}_n\|^2 + d\mathbf{w}^T \mathbf{w}$
- ☐ $\frac{1}{N} \sum_{n=1}^N \|\mathbf{x}_n - \mathbf{w}\mathbf{w}^T \mathbf{x}_n\|^2 + (\mathbf{w}^T \mathbf{w})^2$

Question 6

Nearest Neighbor and RBF Network

Consider getting the 1 Nearest Neighbor hypothesis from a data set of two examples $(\mathbf{x}_+, +1)$ and $(\mathbf{x}_-, -1)$. Which of the following linear hypothesis $g_{LIN}(\mathbf{x}) = \text{sign}(\mathbf{w}^T \mathbf{x} + b)$ (where \mathbf{w} does not include $b = w_0$) is equivalent to the hypothesis?

- ☐ $\mathbf{w} = 2(\mathbf{x}_+ - \mathbf{x}_-), b = -\|\mathbf{x}_+\|^2 + \|\mathbf{x}_-\|^2$
- ☐ $\mathbf{w} = 2(\mathbf{x}_- - \mathbf{x}_+), b = -\mathbf{x}_+^T \mathbf{x}_-$
- ☐ none of the other choices
- ☐ $\mathbf{w} = 2(\mathbf{x}_+ - \mathbf{x}_-), b = +\mathbf{x}_+^T \mathbf{x}_-$
- ☐ $\mathbf{w} = 2(\mathbf{x}_- - \mathbf{x}_+), b = +\|\mathbf{x}_+\|^2 - \|\mathbf{x}_-\|^2$

Question 7

Consider an RBF Network hypothesis for binary classification

$$g_{RBFNET}(\mathbf{x}) = \text{sign}(\beta_+ \exp(-\|\mathbf{x} - \boldsymbol{\mu}_+\|^2) + \beta_- \exp(-\|\mathbf{x} - \boldsymbol{\mu}_-\|^2))$$

and assume that $\beta_+ > 0 > \beta_-$. Which of the following linear hypothesis

$g_{LIN}(\mathbf{x}) = \text{sign}(\mathbf{w}^T \mathbf{x} + b)$ (where \mathbf{w} does not include $b = w_0$) is equivalent to $g_{RBFNET}(\mathbf{x})$?

- ☐ $\mathbf{w} = 2(\boldsymbol{\mu}_+ - \boldsymbol{\mu}_-), b = \ln\left|\frac{\beta_+}{\beta_-}\right| - \|\boldsymbol{\mu}_+\|^2 + \|\boldsymbol{\mu}_-\|^2$
- ☐ $\mathbf{w} = 2(\boldsymbol{\mu}_- - \boldsymbol{\mu}_+), b = \ln\left|\frac{\beta_-}{\beta_+}\right| + \|\boldsymbol{\mu}_+\|^2 - \|\boldsymbol{\mu}_-\|^2$
- ☐ $\mathbf{w} = 2(\beta_+ \boldsymbol{\mu}_+ + \beta_- \boldsymbol{\mu}_-), b = -\beta_+ \|\boldsymbol{\mu}_+\|^2 + \beta_- \|\boldsymbol{\mu}_-\|^2$
- ☐ $\mathbf{w} = 2(\beta_+ \boldsymbol{\mu}_+ + \beta_- \boldsymbol{\mu}_-), b = +\beta_+ \|\boldsymbol{\mu}_+\|^2 - \beta_- \|\boldsymbol{\mu}_-\|^2$
- ☐ none of the other choices

Question 8

Assume that a full RBF network (page 9 of class 214) using $\text{RBF}(\mathbf{x}, \boldsymbol{\mu}) = \exp(-\|\mathbf{x} - \boldsymbol{\mu}\|^2)$ is solved for squared error regression on a data set where all inputs \mathbf{x}_n are different. What are the optimal coefficients β_n for each $\text{RBF}(\mathbf{x}, \mathbf{x}_n)$?

- ☐ $\|\mathbf{x}_n\|^2 y_n^2$
- ☐ $\|\mathbf{x}_n\| y_n$
- ☐ none of the other choices
- ☐ y_n^2
- ☐ y_n

Question 9

Matrix Factorization

Consider matrix factorization of $\tilde{d} = 1$ with alternating least squares. Assume that the $\tilde{d} \times N$ user factor matrix V is initialized to a constant matrix of 1. After step 2.1 of alternating least squares (page 10 of lecture 215), what is the optimal w_m , the $\tilde{d} \times 1$ movie 'vector' for the m -th movie?

- ☐ the average rating of the m -th movie
- ☐ the total rating of the m -th movie
- ☐ the maximum rating of the m -th movie
- ☐ the minimum rating of the m -th movie
- ☐ none of the other choices

Question 10

Assume that for a full rating matrix R , we have obtained a perfect matrix factorization $R = V^T W$. That is, $r_{nm} = \mathbf{v}_n^T \mathbf{w}_m$ for all n, m . Then, a new user ($N + 1$) comes. Because we do not have any information for the type of the movie she likes, we initialize her feature vector \mathbf{v}_{N+1} to $\frac{1}{N} \sum_{n=1}^N \mathbf{v}_n$, the average user feature vector. Now, our system decides to recommend her a movie m with the maximum predicted score $\mathbf{v}_{N+1}^T \mathbf{w}_m$. What would the movie be?

- ☐ none of the other choices
- ☐ the movie with the largest maximum rating
- ☐ the movie with the largest average rating
- ☐ the movie with the largest minimum rating

- ☐ the movie with the smallest rating variance

Question 11

Experiment with Backprop neural Network

Implement the backpropagation algorithm (page 16 of lecture 212) for $d-M-1$ neural network with tanh-type neurons, **including the output neuron**. Use the squared error measure between the output $g_{NNET}(\mathbf{x}_n)$ and the desired y_n and backprop to calculate the per-example gradient.

Because of the different output neuron, your $\delta_1^{(L)}$ would be different from the course slides! Run the algorithm on the following set for training (each row represents a pair of (\mathbf{x}_n, y_n) ; the first column is $(\mathbf{x}_n)_1$; the second one is $(\mathbf{x}_n)_2$; the third one is y_n):

[hw4_nnet_train.dat](#)

and the following set for testing:

[hw4_nnet_test.dat](#)

Fix $T = 50000$ and consider the combinations of the following parameters:

- the number of hidden neurons M
- the elements of $w_{ij}^{(\ell)}$ chosen independently and uniformly from the range $(-r, r)$
- the learning rate η

Fix $\eta = 0.1$ and $r = 0.1$. Then, consider $M \in \{1, 6, 11, 16, 21\}$ and repeat the experiment for 500 times. Which M results in the lowest average E_{out} over 500 experiments?

- ☐ 21
- ☐ 11
- ☐ 6
- ☐ 16
- ☐ 1

Question 12

Following Question 11, fix $\eta = 0.1$ and $M = 3$. Then, consider $r \in \{0, 0.1, 10, 100, 1000\}$ and repeat the experiment for 500 times. Which r results in the lowest average E_{out} over 500 experiments?

- ☐ 0
- ☐ 10
- ☐ 100
- ☐ 1000

☐ 0.1

Question 13

Following Question 11, fix $r = 0.1$ and $M = 3$. Then, consider $\eta \in \{0.001, 0.01, 0.1, 1, 10\}$ and repeat the experiment for 500 times. Which η results in the lowest average E_{out} over 500 experiments?

- ☐ 1
- ☐ 0.001
- ☐ 0.1
- ☐ 10
- ☐ 0.01

Question 14

Following Question 11, deepen your algorithm by making it capable of training a d -8-3-1 neural network with tanh-type neurons. Do not use any pre-training. Let $r = 0.1$ and $\eta = 0.01$ and repeat the experiment for 500 times. Which of the following is true about E_{out} over 500 experiments?

- ☐ $0.06 \leq E_{out} < 0.08$
- ☐ $0.04 \leq E_{out} < 0.06$
- ☐ none of the other choices
- ☐ $0.00 \leq E_{out} < 0.02$
- ☐ $0.02 \leq E_{out} < 0.04$

Question 15

Experiment with 1 Nearest Neighbor

Implement any algorithm that 'returns' the 1\$ Nearest Neighbor hypothesis discussed in page 8 of lecture 214. $g_{\text{nnbor}}(\mathbf{x}) = y_m$ such that \mathbf{x} closest to \mathbf{x}_m

Run the algorithm on the following set for training:

[hw4_knn_train.dat](#)

and the following set for testing:

[hw4_knn_test.dat](#)

Which of the following is closest to $E_{in}(g_{\text{nbor}})$?

- ☐ 0.1
- ☐ 0.0
- ☐ 0.2
- ☐ 0.3
- ☐ 0.4

Question 16

Following Question 15, which of the following is closest to $E_{out}(g_{\text{nbor}})$?

- ☐ 0.34
- ☐ 0.26
- ☐ 0.30
- ☐ 0.32
- ☐ 0.28

Question 17

Now, implement any algorithm for the k Nearest Neighbor with $k = 5$ to get $g_{5\text{-nbor}}(\mathbf{x})$. Run the algorithm on the same sets in Question 15 for training/testing.

Which of the following is closest to $E_{in}(g_{5\text{-nbor}})$?

- ☐ 0.3
- ☐ 0.0
- ☐ 0.1
- ☐ 0.4
- ☐ 0.2

Question 18

Following Question 17, Which of the following is closest to $E_{out}(g_{5\text{-nbor}})$?

- ☐ 0.32
- ☐ 0.26

- ☐ 0.28
- ☐ 0.34
- ☐ 0.30

Question 19

Experiment with k-Means

Implement the k -Means algorithm (page 16 of lecture 214). Randomly select k instances from $\{\mathbf{x}_n\}$ to initialize your $\boldsymbol{\mu}_m$. Run the algorithm on the following set for training:

[hw4_kmeans_train.dat](#)

and repeat the experiment for 500 times. Calculate the clustering E_{in} by

$$\frac{1}{N} \sum_{n=1}^N \sum_{m=1}^M [[\mathbf{x}_n \in S_m]] \|\mathbf{x}_n - \boldsymbol{\mu}_m\|^2$$

as described on page 13 of lecture 214 for $M = k$.

For $k = 2$, which of the following is closest to the average E_{in} of k -Means over 500 experiments?

- ☐ 0.5
- ☐ 2.0
- ☐ 2.5
- ☐ 1.5
- ☐ 1.0

Question 20

For $k = 10$, which of the following is closest to the average E_{in} of k -Means over 500 experiments?

- ☐ 1.5
- ☐ 1.0
- ☐ 0.5
- ☐ 2.0
- ☐ 2.5

☐ In accordance with the Coursera Honor Code, I (Juan Tzu Chuan) certify that the answers here are my own work.

Submit Answers

Save Answers

You cannot submit your work until you agree to the Honor Code. Thanks!